

# Message Passing Interface

# Message Passing Interface (MPI)

- Message Passing Interface (MPI) is a specification designed for parallel applications.
- The goal of the Message Passing Interface is to provide a widely used standard for writing message passing programs. The interface attempts to be
  - practical
  - portable
  - efficient
  - flexible
- A message-passing library specification is:
  - message-passing model
  - not a compiler specification
  - not a specific product

# Message Passing Interface

- MPI is designed to provide access to advanced parallel hardware for
  - end users
  - library writers
  - tool developers
- Message Passing is now mature as programming paradigm
  - well understood
  - efficient match to hardware
  - many applications
- Interface specifications have been defined for C and Fortran programs.

# History of MPI

- April, 1992:
  - Workshop on Standards for Message Passing in a Distributed Memory Environment, sponsored by the Center for Research on Parallel Computing, Williamsburg, Virginia.
  - The basic features essential to a standard message passing interface were discussed, and a working group established to continue the standardization process. Preliminary draft proposal developed subsequently.
- November 1992:
  - Working group meets in Minneapolis. Oak Ridge National Laboratory (ORNL) presented a MPI draft proposal (MPI1).
  - Group adopts procedures and organization to form the [MPI Forum](#). MPIF eventually comprised of about 175 individuals from 40 organizations including parallel computer vendors, software writers, academia and application scientists.

# History of MPI

- November 1993: Supercomputing 93 conference - draft MPI standard presented.
- Final version of draft released in May, 1994 - available on the WWW at: <http://www.mcs.anl.gov/Projects/mpi/standard.html>
- MPI-2 picked up where the first MPI specification left off, and addresses topics which go beyond the first MPI specification.

# Reasons for Using MPI

- **Standardization** - MPI is the only message passing library which can be considered a standard. It is supported on virtually all HPC platforms.
- **Portability** - There is no need to modify your source code when you port your application to a different platform that supports (and is compliant with) the MPI standard.
- **Performance Opportunities** - Vendor implementations should be able to exploit native hardware features to optimize performance.
- **Functionality** - Over 115 routines are defined.
- **Availability** - A variety of implementations are available, both vendor and public domain.

# Types of Parallel Programming

- Flynn's taxonomy (hardware oriented)
  - SISD : Single Instruction, Single Data
  - SIMD : Single Instruction, Multiple Data
  - MISD : Multiple Instruction, Single Data
  - MIMD : Multiple Instruction, Multiple Data
- A programmer-oriented taxonomy
  - Data-parallel : Same operations on different data (SIMD)
  - Task-parallel : Different programs, different data
  - MIMD : Different programs, different data
  - SPMD : Same program, different data
  - Dataflow : Pipelined parallelism
- MPI is for SPMD/MIMD.

# MPI Programming

- Writing MPI programs
- Compiling and linking
- Running MPI programs
- More information
  - *Using MPI* by William Gropp, Ewing Lusk, and Anthony Skjellum,
  - *Designing and Building Parallel Programs* by Ian Foster.
  - A Tutorial/User's Guide for MPI by Peter Pacheco (ftp://math.usfca.edu/pub/MPI/mpi.guide.ps)
  - The MPI standard and other information is available at <http://www.mcs.anl.gov/mpi>. Also the source for several implementations.



# Simple MPI C Program (hello.c)

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char **argv )
{
    MPI_Init(&argc, &argv);
    printf("Hello world\n");
    MPI_Finalize();
    return 0;
}
```

- Compile hello.c: `mpicc -o hello hello.c`
- `#include <mpi.h>` provides basic MPI definitions and types
- `MPI_Init` starts MPI
- `MPI_Finalize` exits MPI
- Note that all non-MPI routines are local; thus the `printf` run on each process

# Simple MPI C++ Program (hello.cc)

```
#include <iostream.h>
```

```
#include <mpi.h>
```

```
Int main(int argc, char *argv[])
```

```
{
```

```
    int rank, size;
```

```
    MPI::Init(argc, argv);
```

```
    cout << "Hello world" << endl;
```

```
    MPI::Finalize();
```

```
    return 0;
```

```
}
```

- Compile hello.cc: `mpiCC -o hello hello.cc`

# Starting and Exiting MPI

- Starting MPI

```
int MPI_Init(int *argc, char **argv)
```

```
void MPI::Init(int& argc, char**& argv)
```

- Exiting MPI

```
int MPI_Finalize(void)
```

```
void MPI::Finalize()
```

# Running MPI Programs

- On many platforms MPI programs can be started with ‘`mpirun`’.

```
mpirun -np <np> hello
```

- ‘`mpirun`’ is not part of the standard, but some version of it is common with several MPI implementations.
- Two of the first questions asked in a parallel program are as follows:
  1. “How many processes are there?”
  2. “Who am I?”
- “How many” is answered with `MPI_COMM_SIZE`;  
“Who am I” is answered with `MPI_COMM_RANK`.
- The rank is a number between zero and (`SIZE - 1`).

# Second MPI C Program

```
#include <stdio.h>
#include "mpi.h"

int main(int argc, char **argv)
{
    int rank, size;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    printf("Hello world! I am %d of %d\n",rank,size);
    MPI_Finalize();
    return 0;
}
```

# Second MPI C++ Program

```
#include <iostream.h>
```

```
#include <mpi.h>
```

```
int main(int argc, char **argv)
```

```
{
```

```
    MPI::Init(argc, argv);
```

```
    int rank = MPI::COMM_WORLD.Get_rank();
```

```
    int size = MPI::COMM_WORLD.Get_size();
```

```
    cout << "Hello world! I am " << rank << " of " << size << endl;
```

```
    MPI::Finalize();
```

```
    return 0;
```

```
}
```

# Second MPI C Program

```
#include <stdio.h>
#include "mpi.h"

int main(int argc, char **argv)
{
    int rank, size;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    printf("Hello world! I am %d of %d\n",rank,size);
    MPI_Finalize();
    return 0;
}
```

# MPI\_COMM\_WORLD

- Communication in MPI takes place with respect to communicators (more about communicators later).
- The MPI\_COMM\_WORLD communicator is created when MPI is started and contains all MPI processes.
- MPI\_COMM\_WORLD is a useful default communicator – many applications do not need to use any other.
- `mpirun -np <number of processes> <program name and arguments>`



# MPI C Data Type

MPI C Data Type	C Data Type
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED_INT	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	
MPI_PACKED	

# MPI C++ Data Types

MPI C++ data type	C++ data type
MPI::CHAR	signed char
MPI::SHORT	signed short int
MPI::INT	signed int
MPI::LONG	signed long int
MPI::UNSIGNEDCHAR	unsigned char
MPI::UNSIGNEDSHORT	unsigned short int
MPI::UNSIGNED	unsigned int
MPI::UNSIGNEDLONG	unsigned long int
MPI::FLOAT	float
MPI::DOUBLE	double
MPI::LONGBDOUBLE	long double
MPI::BYTE	
MPI::PACKED	

# MPI Basic Functions

- **MPI\_Init** – set up an MPI program.
- **MPI\_Comm\_size** – get the number of processes participating in the program.
- **MPI\_Comm\_rank** – determine which of those processes corresponds to the one calling the function.
- **MPI\_Send** – send messages.
- **MPI\_Recv** – receive messages.
- **MPI\_Finalize** – stop participating in a parallel program.

# MPI Basic Functions

- **MPI\_Init(int \*argc, char \*\*\*argv)**
  - Takes the command line arguments to a program,
  - checks for any MPI options, and
  - passes remaining command line arguments to the main program.
- **MPI\_Comm\_size( MPI\_Comm comm, int \*size )**
  - Determines the size of a given MPI Communicator.
  - A communicator is a set of processes that work together.
  - For typical programs this is the default MPI\_COMM\_WORLD, which is the communicator for all processes available to an MPI program.

# MPI Basic Functions

- **MPI\_Comm\_rank( MPI\_Comm comm, int \*rank )**
  - Determine the rank of the current process within a communicator.
  - Typically, if a MPI program is being run on N processes, the communicator would be MPI\_COMM\_WORLD, and the rank would be an integer from 0 to N-1.
- **MPI\_Send( void \*buf, int count, MPI\_Datatype datatype, int dest, int tag, MPI\_Comm comm )**
  - Send the contents of **buf**, which contains count elements of type **datatype** to a process of rank **dest** in the communicator **comm**, flagged with the message **tag**. Typically, the communicator is MPI\_COMM\_WORLD.

# MPI Basic Functions

- **MPI\_Recv( void \*buf, int count, MPI\_Datatype datatype, int source, int tag, MPI\_Comm comm, MPI\_Status \*status )**
  - Read into **buf**, which contains **count** elements of type **datatype** from process **source** in communicator **comm** if a message is sent flagged with **tag**. Also receive information about the transfer into **status**.
- **MPI\_Finalize()**
  - Handles anything that the current MPI protocol will need to do before exiting a program.
  - Typically should be the final or near final line of a program.

# MPI Functions

- **MPI\_Bcast** – Broadcasts a message from the process with rank "root" to all other processes of the group.
  - **int MPI\_Bcast ( void \*buffer, int count, MPI\_Datatype datatype, int root, MPI\_Comm comm)**
  - Input/output Parameters
    - **buffer** starting address of buffer (choice)
    - **count** number of entries in buffer (integer)
    - **datatype** data type of buffer (handle)
    - **root** rank of broadcast root (integer)
    - **comm** communicator (handle)

# MPI Functions

- **MPI\_Reduce** – Reduces values on all processes to a single value
  - `int MPI_Reduce ( void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm)`
  - Input/output Parameters
    - **sendbuf** address of send buffer (choice)
    - **count** number of elements in send buffer (integer)
    - **datatype** data type of elements of send buffer (handle)
    - **op** reduce operation (handle)
    - **root** rank of root process (integer)
    - **comm** communicator (handle)



# MPI Reduce Operation

MPI Name	Operation
MPI_MAX	Maximum
MPI_MIN	Minimum
MPI_PROD	Product
MPI_SUM	Sum
MPI_LAND	Logical and
MPI_LOR	Logical or
MPI_LXOR	Logical exclusive or (xor)
MPI_BAND	Bitwise and
MPI_BOR	Bitwise or
MPI_BXOR	Bitwise xor
MPI_MAXLOC	Maximum value and location
MPI_MINLOC	Minimum value and location

# MPI Example Programs

- **Master/Worker** (greeting.c, master-worker.c, ring.c) – A (master) process waits for the message sent by other (worker) processes.
- **Integration** (integrate.c): Evaluates the trapezoidal rule estimate for an integral of  $F(x)$ .
- **$\pi$  Calculation** (pi.c) – Each process calculates part of  $\pi$  and summarize in the master process.
- **Matrix Multiplication** (matrix-multiply.c, matrix.c) – Each worker process calculates some rows of matrix multiplication and sends the result back to the master process.

# MPI Example Programs

- **Exponential** (exponent.c): Evaluates  $e$  for a series of cases. This is essentially the sum of products of  $(1 - \exp())$ .
- **Statistics** (statistics.c) : Tabulates statistics of a set of test scores. :
- **Tridiagonal system** (tridiagonal.c) : Solves the tridiagonal system  $Tx = y$ .